

Pollack's Rule

Justification for Heterogeneous Computing

Big picture

- Performance modeling: Estimating performance of a hypothetical system so system designer can compare different options
- Today: Consider different configurations of cores
 - Assumption: Total processor size (silicon area) is the same for all configurations

Pollack's rule

- The performance of a processing core is proportional to the square root of its area

Pollack's rule

- The performance of a processing core is proportional to the square root of its area

If a single core is replaced by 4 cores, each $\frac{1}{4}$ as large, what is the expected peak performance of the entire system? (i.e. the performance assuming all 4 could be kept perfectly busy)

- A. Half as much as before
- B. The same as before
- C. Twice as much as before
- D. Four times as much as before
- E. None of the above

Pollack's rule

- The performance of a processing core is proportional to the square root of its area

If a single core is replaced by 4 cores, each $\frac{1}{4}$ as large, what is the expected peak performance of the entire system? (i.e. the performance assuming all 4 could be kept perfectly busy)

- A. Half as much as before
- B. The same as before
- C. Twice as much as before
- D. Four times as much as before
- E. None of the above

How does the running time change when a single core is replaced with 4 cores if only half the program can be parallelized?

- Parallel part:

$$\frac{1}{2} \text{ the work} / 2 \text{ the performance} = \frac{1}{4}$$

How does the running time change when a single core is replaced with 4 cores if only half the program can be parallelized?

- Parallel part:

$$\frac{1}{2} \text{ the work} / 2 \text{ the performance} = \frac{1}{4}$$

- Serial part:

$$\frac{1}{2} \text{ the work} / \frac{1}{2} \text{ the performance} = 1$$

Total time: 1.25 times as long

Recall: Amdahl's Law

$$T_p = \underbrace{\frac{T_1(1-B)}{p}}_{\text{Time for parallel part}} + \underbrace{T_1 B}_{\text{Time for serial part}}$$

T_p = processing time on p processors

T_1 = processing time on 1 processor

B = fraction of program that can run in parallel

By what factor does the running time of a program that can be 75% parallelized change on 4 equal-sized cores?

- A. 0.6
- B. 0.875
- C. 1
- D. 1.35
- E. None of the above

By what factor does the running time of a program that can be 75% parallelized change on 4 equal-sized cores?

- A. 0.6
- B. 0.875 $(0.75/2 + 0.25/0.5)$
- C. 1
- D. 1.35
- E. None of the above

By what factor does the running time of a program that can be 90% parallelized change on 4 equal-sized cores?

- A. 0.45
- B. 0.65
- C. 0.765
- D. 1
- E. None of the above

By what factor does the running time of a program that can be 90% parallelized change on 4 equal-sized cores?

- A. 0.45
- B. 0.65 ($0.9/2 + 0.1/0.5$)
- C. 0.765
- D. 1
- E. None of the above


Factor by which running time changes for different programs

% of program that can be parallelized


	75%	90%	95%
1 core	1	1	1
4 cores	0.875	0.65	0.575
9 cores	1	0.6	0.467
16 cores	1.1875	0.625	0.438
25 cores	1.4	0.68	0.44
36 cores	1.625	0.75	0.458

As the number of cores increases, highly parallelizable programs have improved performance, but less parallelizable programs suffer

What about unequal core sizes?




What about unequal core sizes?



By what factor does the peak performance of this system differ from a single core?

- A. ≈ 0.707
- B. ≈ 1.207
- C. ≈ 1.707
- D. ≈ 2.121
- E. None of the above


What about unequal core sizes?



By what factor does the peak performance of this system differ from a single core?

- A. ≈ 0.707
- B. ≈ 1.207
- C. ≈ 1.707
- D. ≈ 2.121
- E. None of the above


What about unequal core sizes?



By what factor does the running time of a program that can be 75% parallelized change?

- A. ≈ 0.793
- B. ≈ 0.854
- C. ≈ 0.939
- D. ≈ 1
- E. None of the above

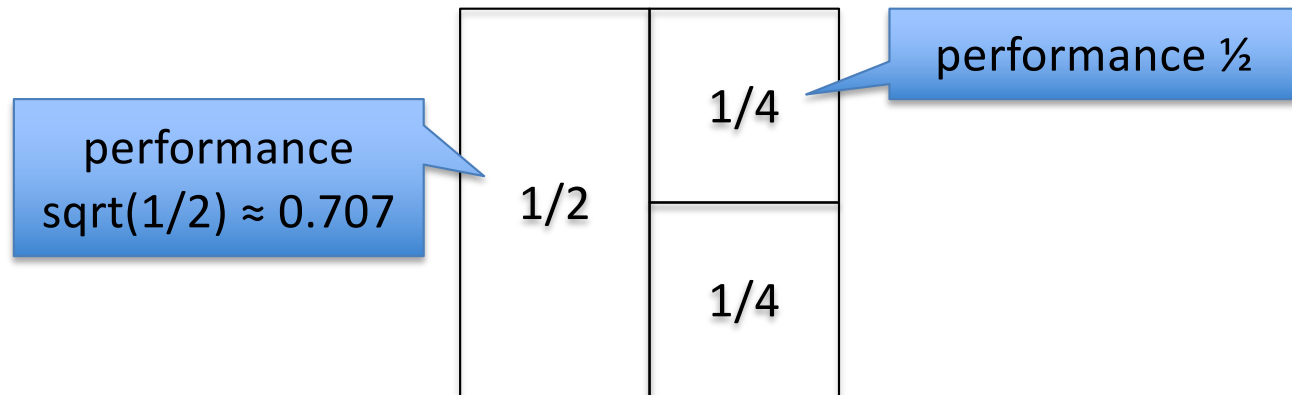
What about unequal core sizes?



By what factor does the running time of a program that can be 75% parallelized change?

- A. ≈ 0.793 ($0.75/1.707 + 0.25/0.707$)
- B. ≈ 0.854
- C. ≈ 0.939
- D. ≈ 1
- E. None of the above


What about unequal core sizes?



By what factor does the running time of a program that cannot be parallelized change?

- A. ≈ 0.707
- B. ≈ 1
- C. ≈ 1.207
- D. ≈ 1.414
- E. None of the above

What about unequal core sizes?



By what factor does the running time of a program that cannot be parallelized change?

- A. ≈ 0.707
- B. ≈ 1
- C. ≈ 1.207
- D. ≈ 1.414 $(1/0.707)$
- E. None of the above

Factor by which running time changes for different programs

	% of program that can be parallelized		
	50%	75%	90%
4 equal cores	1.25	0.88	0.65
Half-sized + 2 quarter-sized cores	1.00	0.79	0.66

Having different sized cores improves performance on less parallelizable programs at small cost on more highly parallelizable ones

Heterogeneity on a cell phone

The screenshot shows the 'Benchmarks' app interface on an Android phone. The top status bar displays the time 9:56, signal strength, and 84% battery. The app's header is blue with a hamburger menu icon and the title 'Benchmarks'. Below the header are three tabs: 'CPU', 'COMPUTE', and 'BATTERY'. The 'CPU' tab is selected. The main content area is titled 'YOUR DEVICE' and lists the following specifications:

Model	OnePlus 5T
OS	Android 9
CPU	Qualcomm MSM8998 Snapdragon 835
Cluster 1	4 Cores @ 1.90 GHz
Cluster 2	4 Cores @ 2.46 GHz

The last two rows of the table are highlighted with a green border. Below the device specifications is a section titled 'CPU BENCHMARK' with a descriptive paragraph and a 'RUN CPU BENCHMARK' button.

8 cores, 2 levels of performance